

Apache Kafka Reinvented for the Cloud

Loved by Developers, Trusted by Enterprises

Contents

- 3 Introduction
- 6 **Cloud-Native:** How a truly cloud-native Kafka solution can modernize your architecture
- 10 **Complete:** How a complete and secure Kafka solution can improve time-to-value and TCO
- 13 **Everywhere:** Build hybrid and multi-cloud architectures with a platform that exists everywhere your data resides
- 15 Customer stories
- 16 Why you need Confluent



Introduction

Apache Kafka® has gone mainstream. One of the most popular open source projects, Kafka is used by over 70% of the Fortune 500 to modernize their data architectures and power event-driven applications that deliver rich, digital customer experiences and data-driven backend operations. Kafka has become the central nervous system for many enterprises, enabling them to connect, process, and react to their data in real time across all of their environments.

However, while Kafka is core to the modern data stack, it doesn't offer everything needed out-of-the-box. Connectors need to be developed and maintained, security features for access control & encryption custom-built, and disaster recovery tools & failover logic designed. And once that's completed, the platform creates a significant ongoing operational burden to maintain its performance, security, and availability going forward to protect against business-disrupting downtime or data loss.

What does this all mean for your business? Higher total cost of ownership, delayed time-to-value, and lower ROI on your mission-critical data streaming projects.

Confluent provides everything you need to implement Kafka use cases quickly, safely, and reliably, so you can instead focus on developing real-time applications that drive the real value for your business. To do so, you need a data streaming platform that frees your organization from spending valuable development cycles on building and maintaining foundational tooling. Your teams have better things to do.

HERE'S HOW CONFLUENT CAN HELP Confluent offers a truly cloud-native Apache Kafka experience with a complete set of enterprise-grade features, available everywhere your data resides.

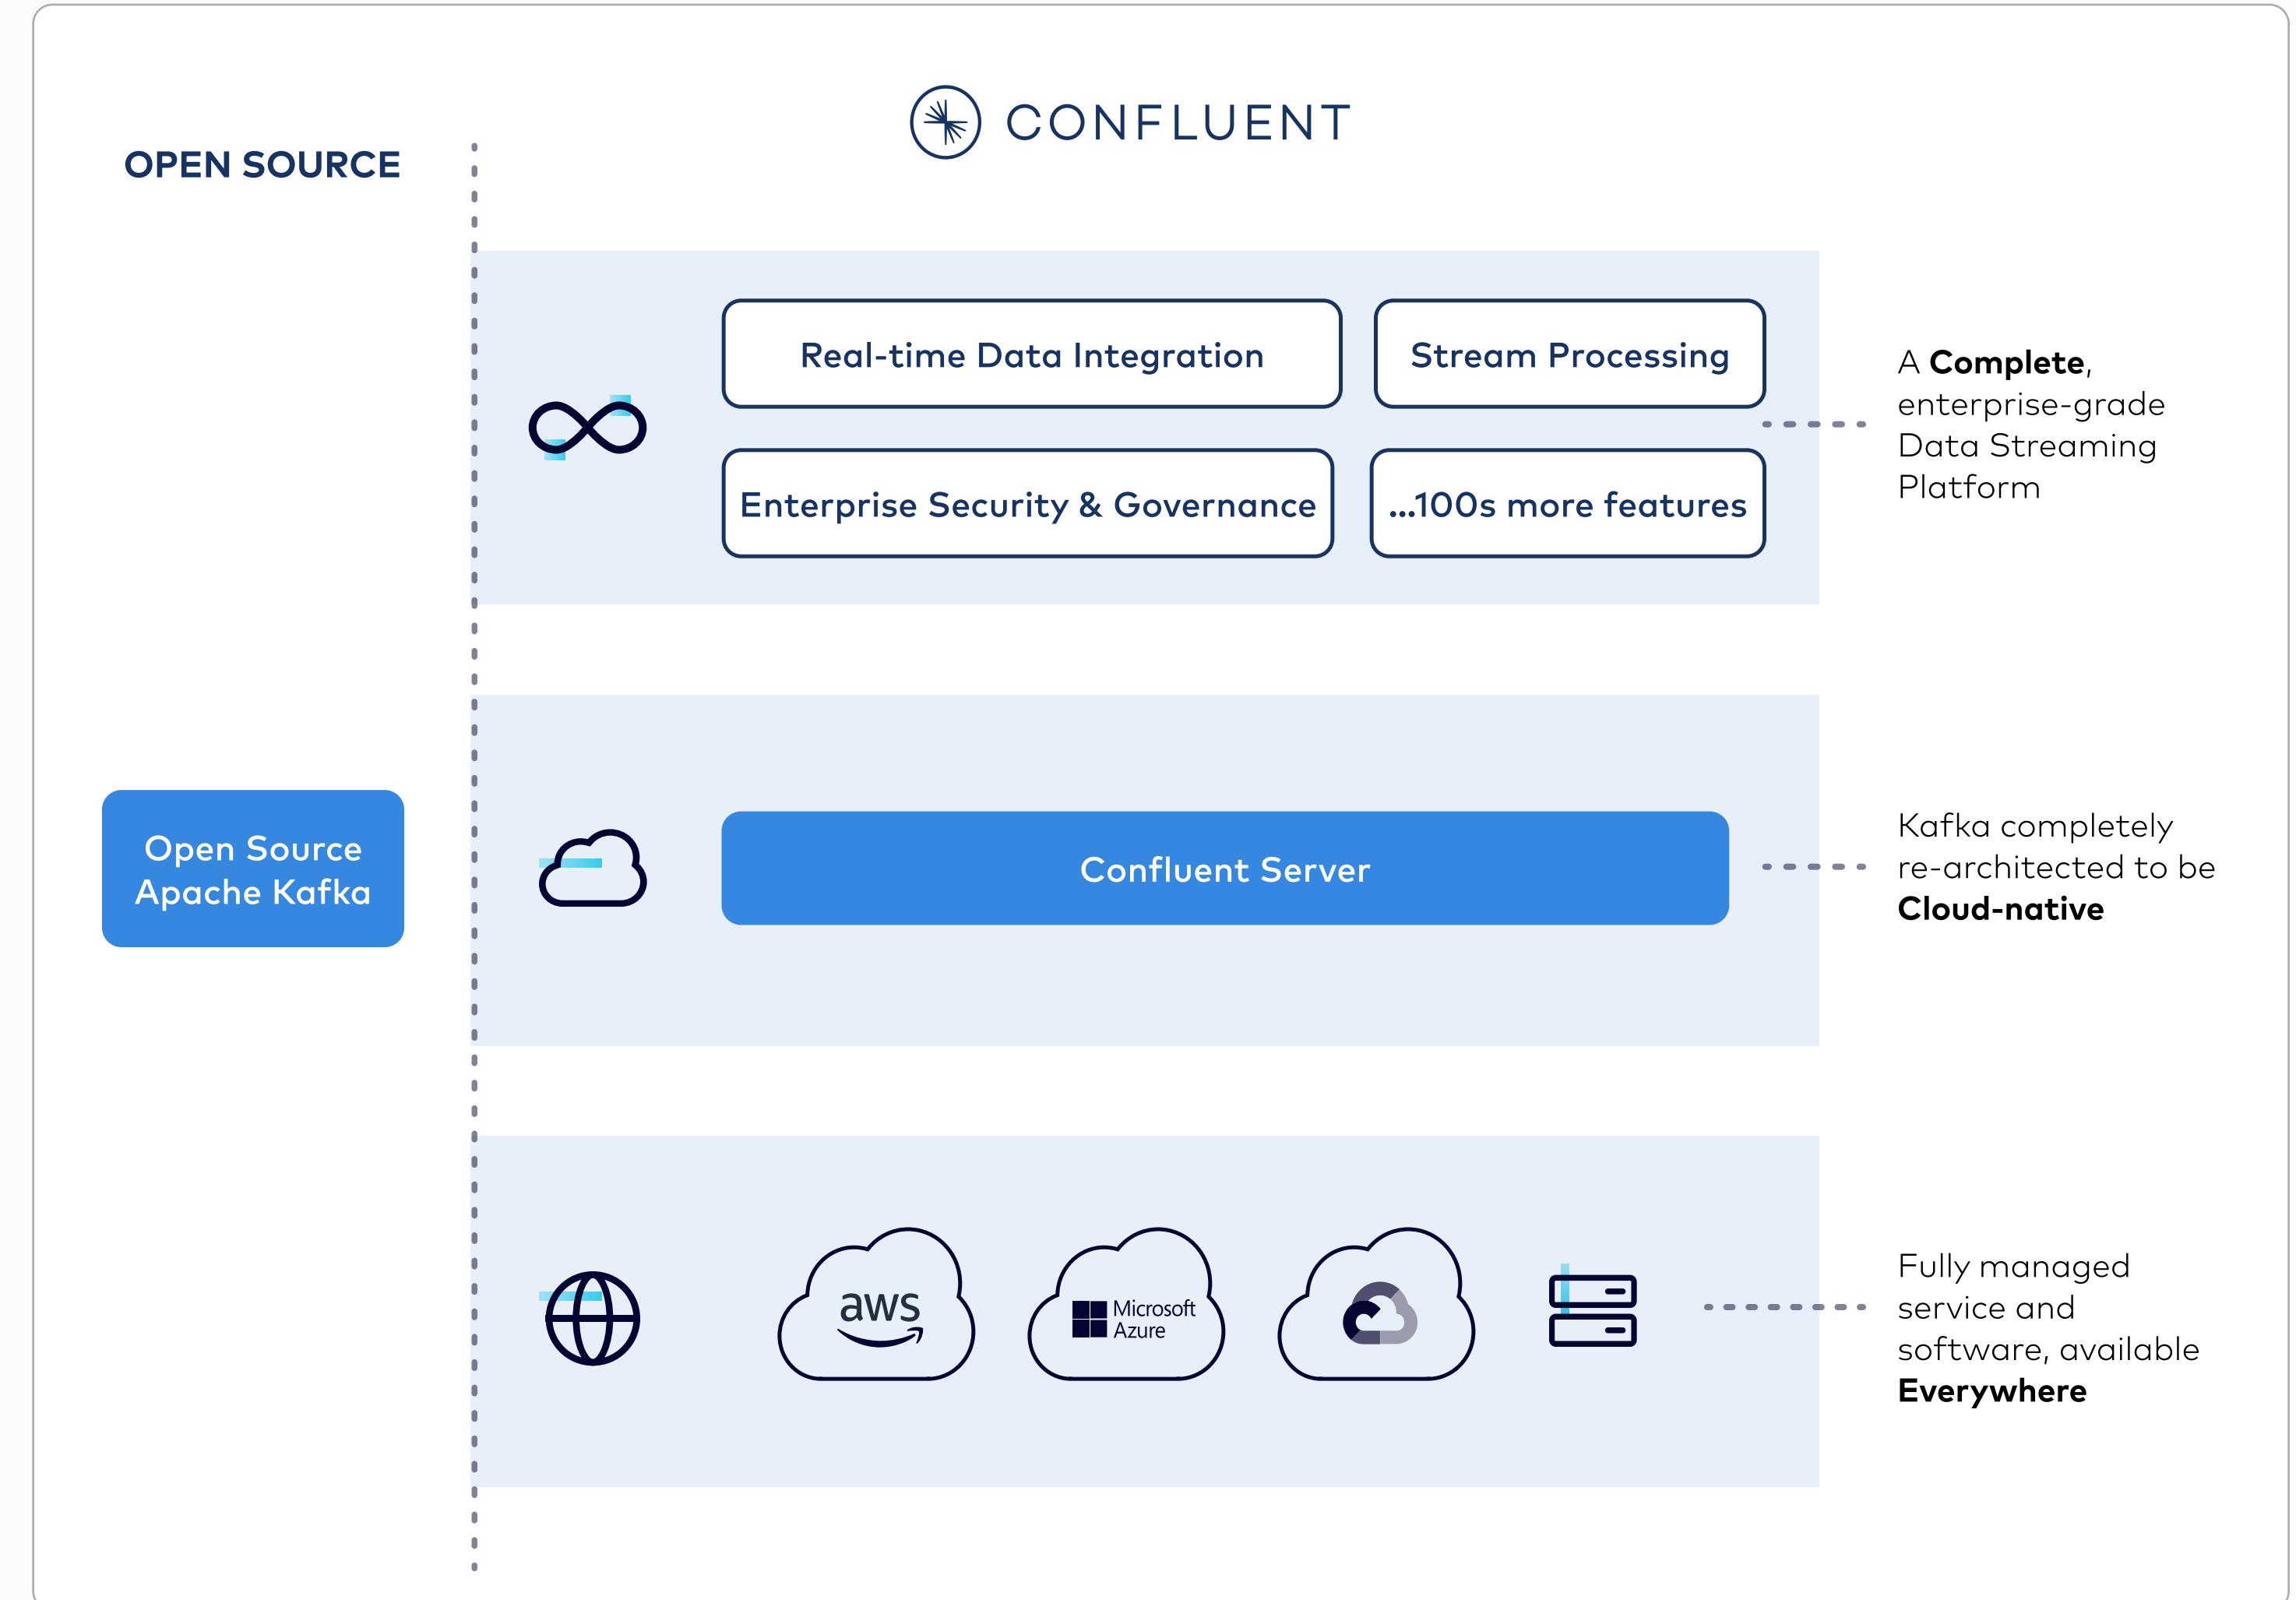




Kafka Reinvented for the Data Streaming Era

Confluent has invested 5M+ engineering hours in completely re-architecting Kafka to harness the full power of the cloud and support all of your mission-critical enterprise workloads. There are three attributes that set Confluent apart from Apache Kafka and other solutions for data streaming:

- Confluent is **cloud-native**, providing a 10X better, fully managed service to free you from operational burdens and risks
- Confluent is **complete** with all the essential tools for developers to build quickly, reliably, and securely
- Confluent exists **everywhere** to connect all of your data in real time from on-prem to cloud and across clouds





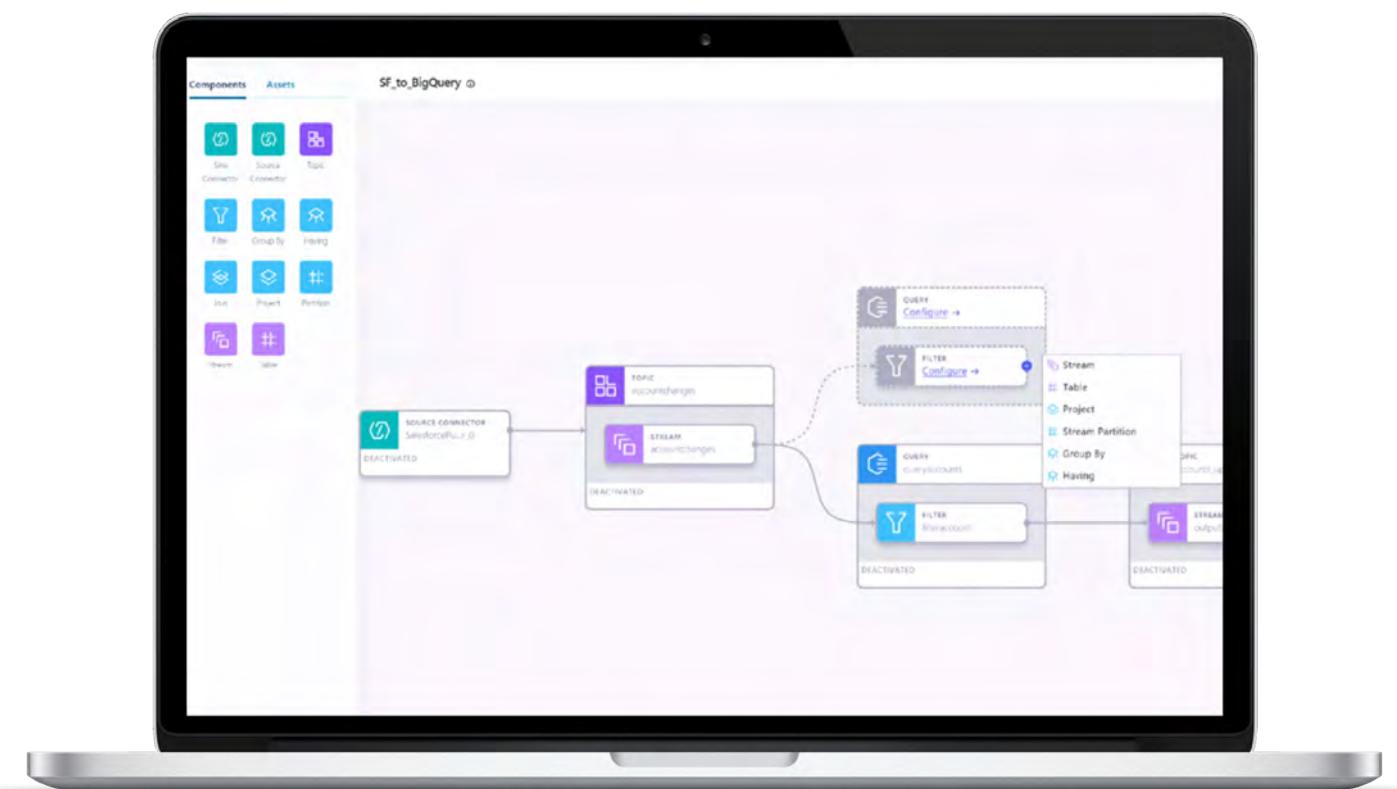
Cloud-native

Confluent has completely re-engineered Kafka to be a truly cloud-native data streaming platform, with 10x faster scaling, infinitely more storage, and built-in resilience. With Confluent Cloud, you can deploy, operate, and scale your Kafka projects in minutes instead of months—all without the overhead and headaches of self-management. Confluent can instantly provision clusters that expand or shrink elastically to meet fluctuating demand, enable cost-effective data storage at any scale, and ensure your clusters are highly available and resilient with our 99.99% SLA.

And for organizations still early in their journey to the cloud, Confluent Platform provides many of the same cloud-native attributes with the control and customization that comes with self-managing infrastructure in your private environments.

Complete

Confluent completes Apache Kafka with a portfolio of essential, enterprise-grade capabilities to securely and efficiently power all of your mission critical streaming workloads. We provide a set of features ready out-of-the-box to unleash developer productivity, operate efficiently at scale, and meet stringent security and availability requirements before moving into production.



Everywhere

Confluent's solution doesn't cover just a subset of cloud providers or on-prem environments. Instead, Confluent offers true deployment flexibility by meeting you everywhere your data resides, from on-premises to hybrid cloud to multi-cloud.

For workloads in the public cloud, Confluent offers a fully managed service available on AWS, Azure, and Google Cloud. For workloads in your private cloud or traditional data centers, Confluent offers a self-managed, enterprise-grade distribution of Kafka that secures and automates cluster deployment and key lifecycle tasks across bare metal servers, virtual machines, and Kubernetes environments.

Confluent can also support your hybrid and multi-cloud strategies by easily linking clusters across different environments.



Our offerings also include partnership for business success, ensuring your adoption of the transformational platform is seamless and creates significant value for your organization. Underpinning everything is our committer-driven expertise, providing support, professional services, and training from the team with over 5 million hours of experience with Apache Kafka.

01

Cloud-Native

How a truly cloud-native Kafka solution can modernize your architecture

DATA SYSTEMS BUILT FOR THE CLOUD have significant advantages in efficiency, speed, and flexibility, all while offering an elevated user experience. Confluent re-architected the Kafka experience for the cloud, with 10x faster scaling, infinite storage, and built-in resiliency.





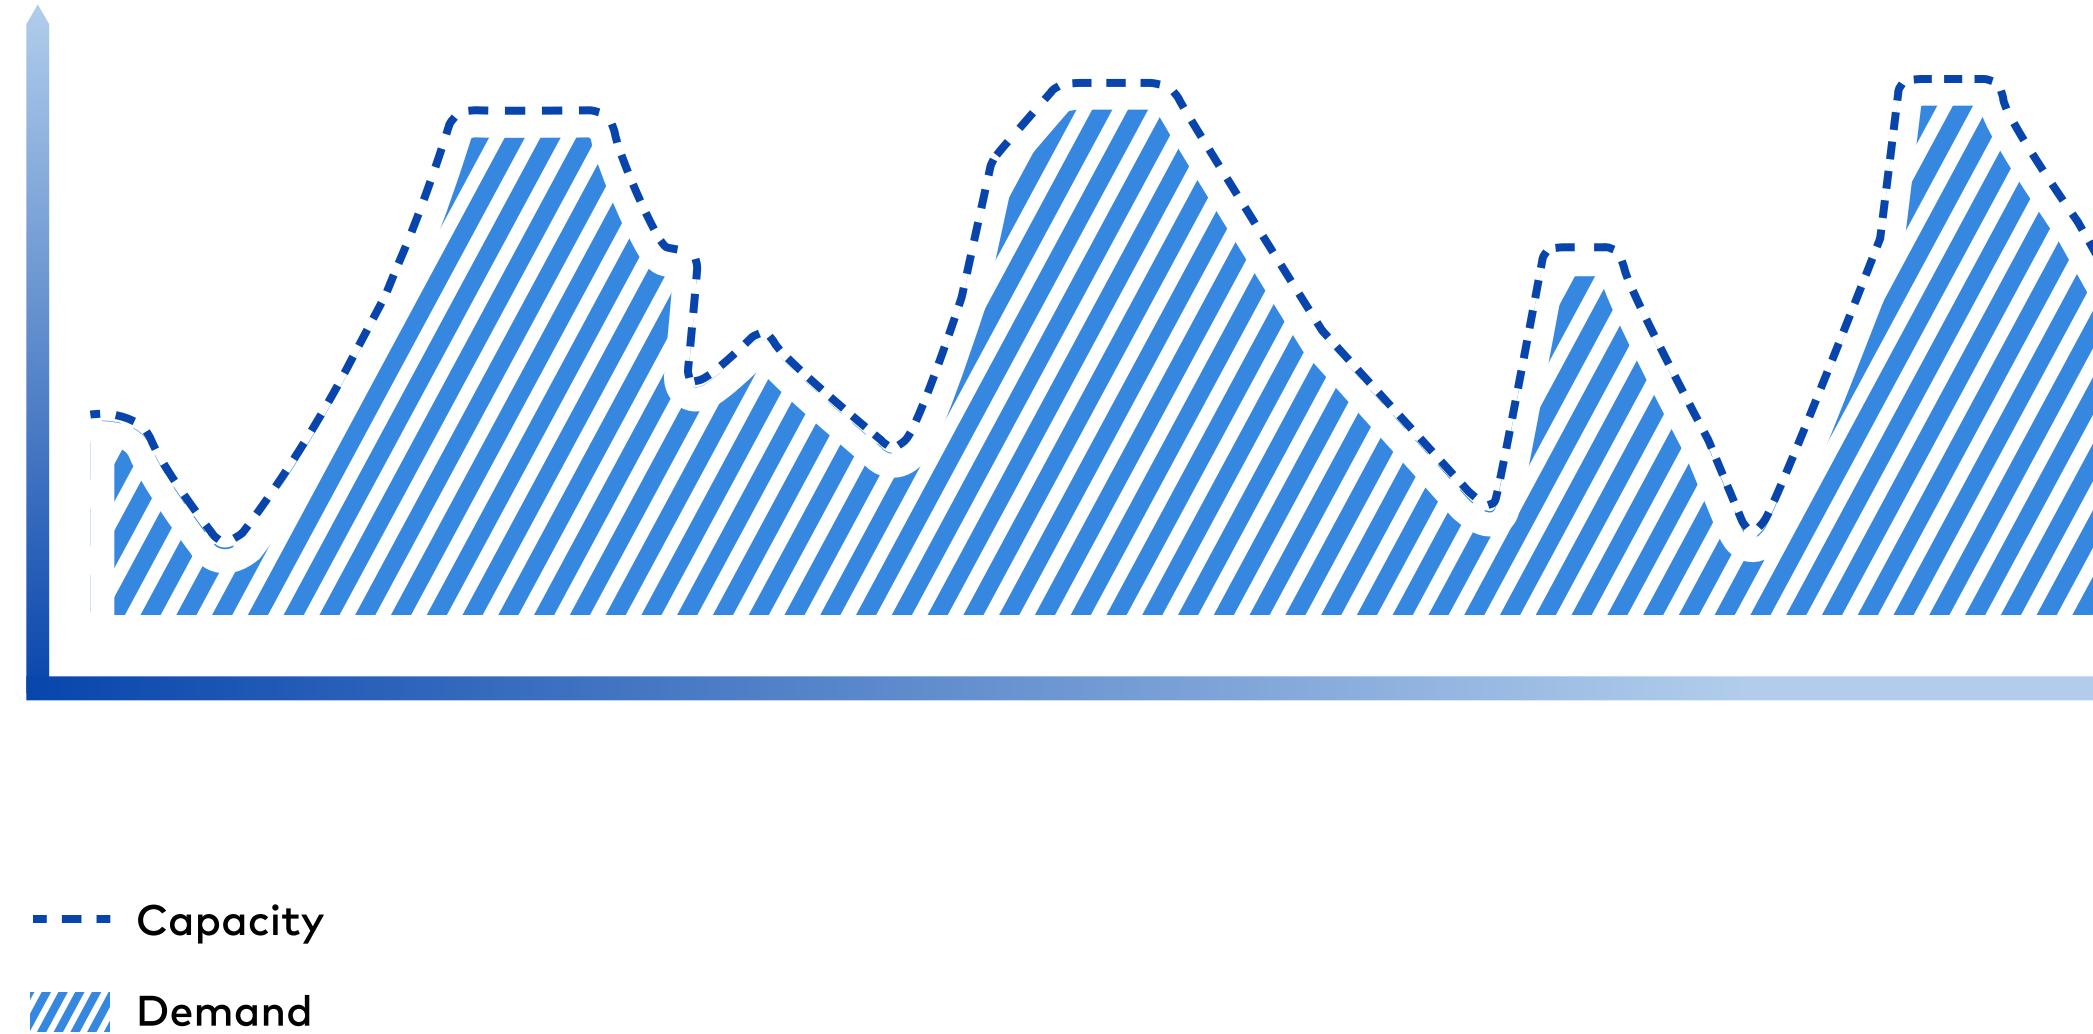
10x Elasticity: Instantly scale your clusters to meet any demand

A key advantage of cloud-native data systems is the ability to elastically scale up as demand peaks, and seamlessly scale down as demand falls. Keeping your Kafka deployments afloat is expensive, time consuming, and risky—especially as usage scales and your mission-critical services need to maintain high availability at all times. Whether it's handling traffic spikes from unpredictable current events, social trends, seasonality, or even the holidays, your teams are often forced to overprovision unnecessarily to ensure capacity and uptime. Confluent dramatically simplifies the capacity planning and scaling processes for you. With Confluent, your teams can:

- **Instantly meet demand** by handling any workload spikes
- **Avoid over-provisioning** with elastic expanding and shrinking of Kafka clusters
- **Eliminate burdens** of complex capacity planning and scaling processes

Elastic: Efficiently scale clusters to meet any demand

Throughput / Data Traffic Over Time





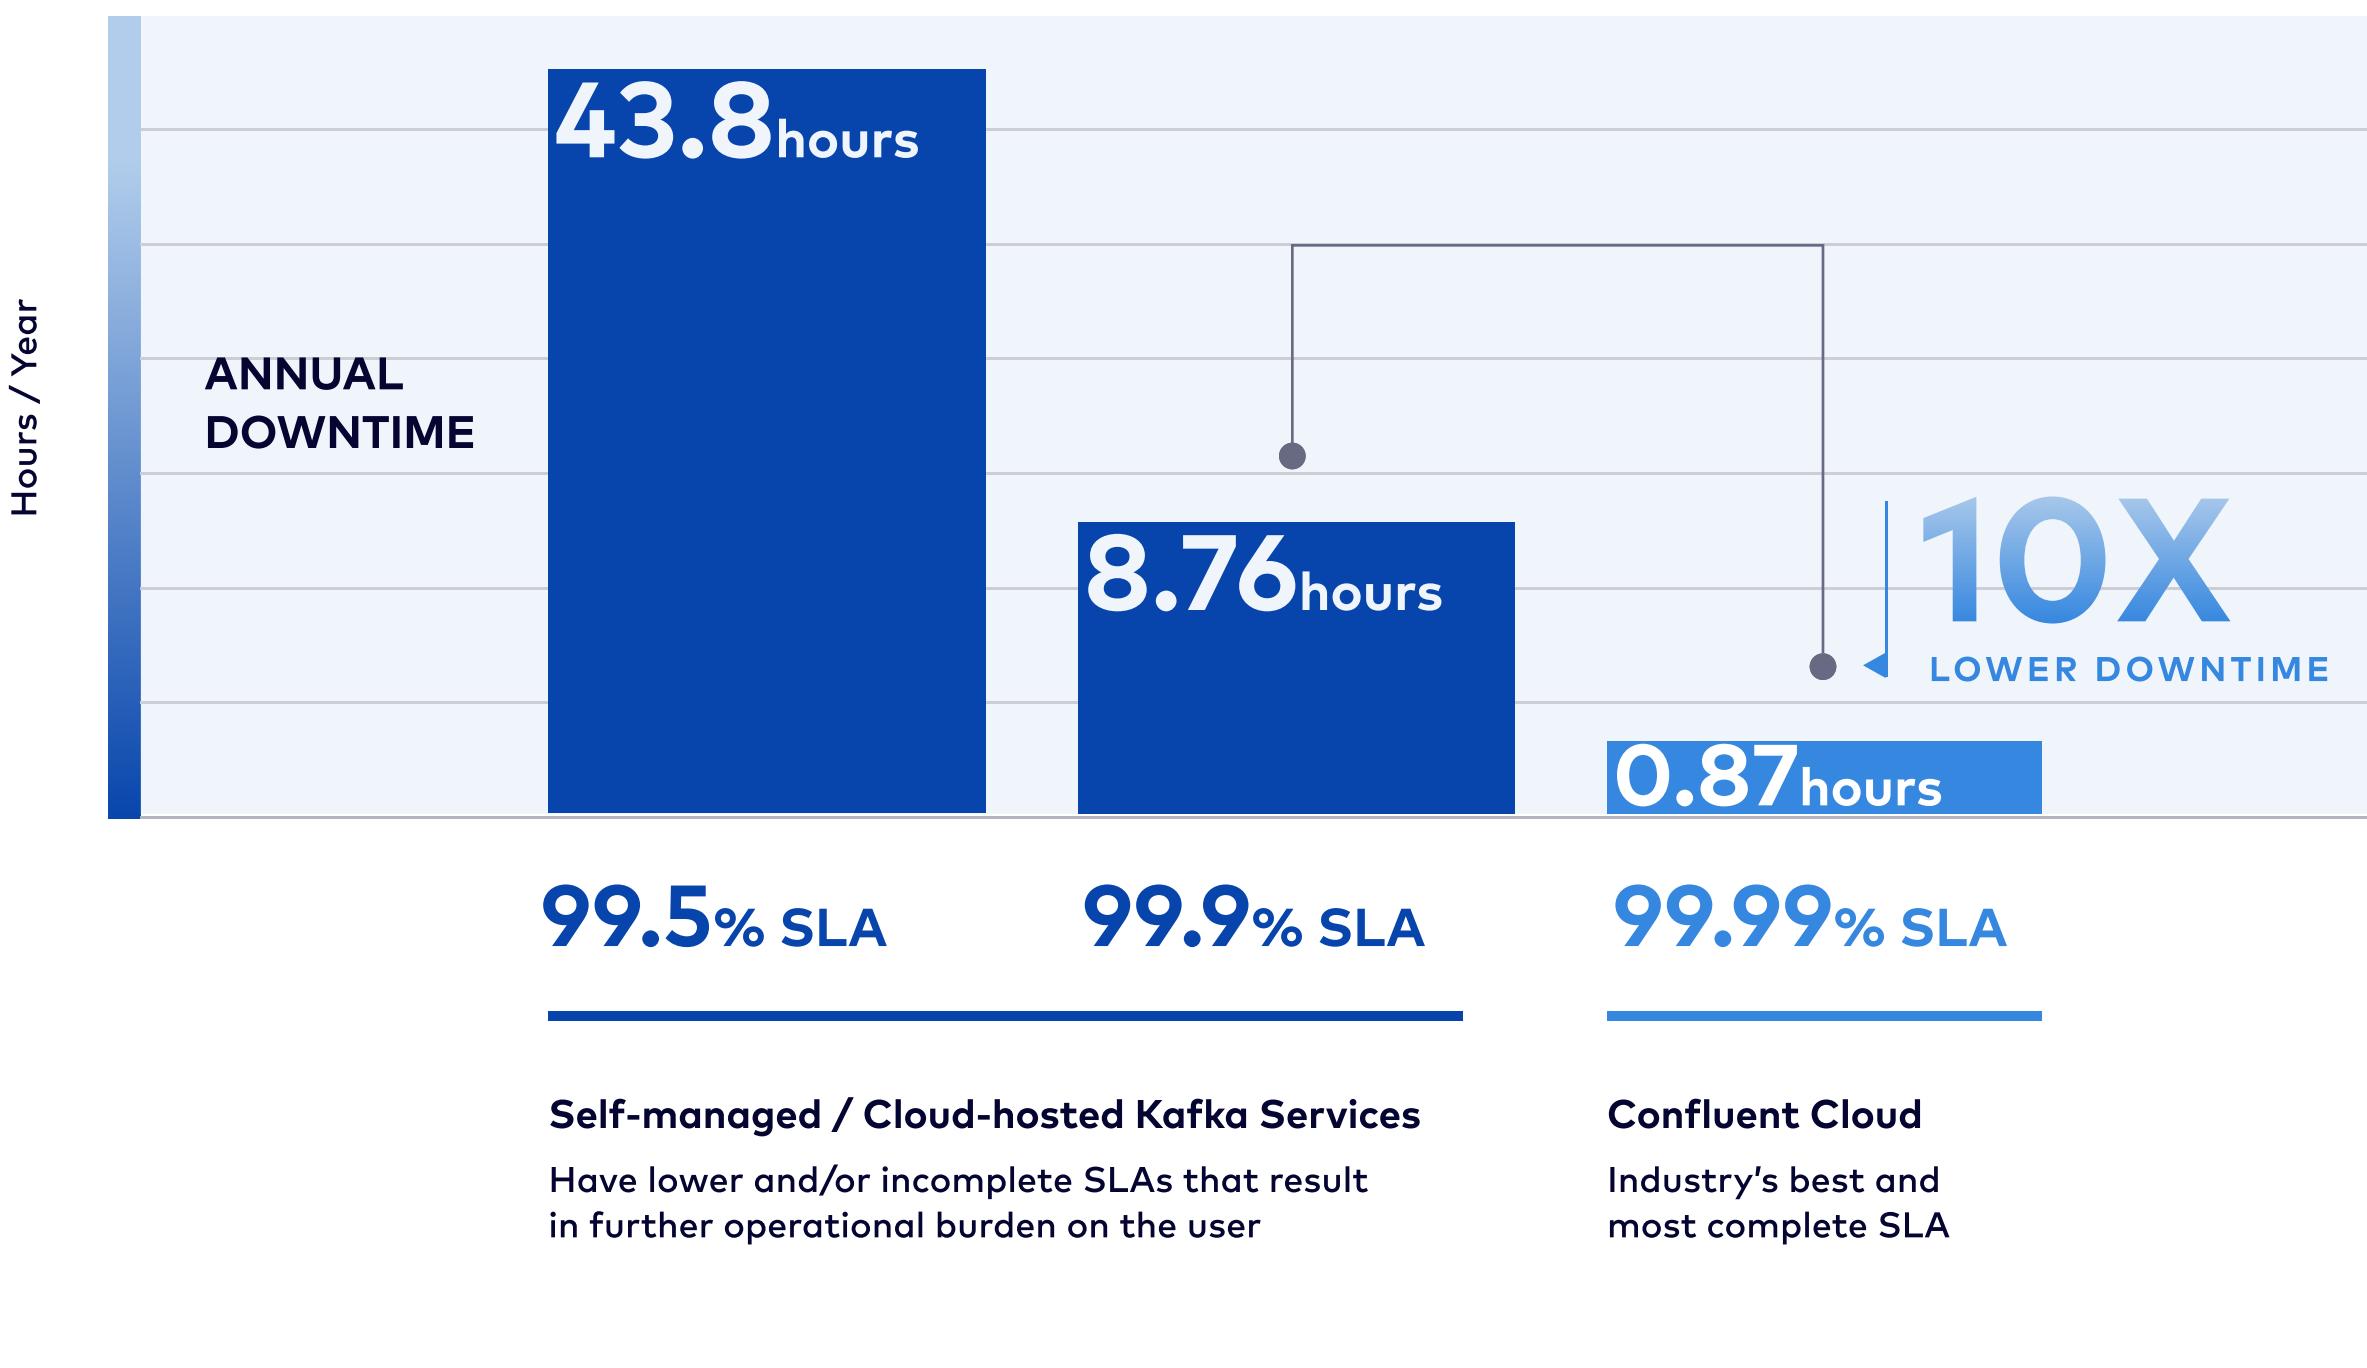
10x Resiliency: Never worry about Kafka-related downtime again

Kafka is a powerful distributed system, but it can be complex to keep clusters running at peak performance and ensure high availability. It's also very challenging to preemptively identify and solve problems that lead to lag, downtime, or data loss. The resiliency of your data streaming platform is non-negotiable, and solely relying on open source Kafka's underlying design is not always enough.

With Confluent, you can leave all your reliability worries behind with the peace of mind of our 10X better availability and durability.

- **Minimize downtime** with our 99.99% uptime SLA without the associated operational burden and overhead
- **Avoid data loss or corruption** with proactive and robust real-time auditing services
- **Offload Kafka maintenance** to a fully managed service that handles all your deployment, optimization, upgrades, and recovery for you

The smallest difference in SLA means exponentially more downtime for your customers





10x Storage: Cost-effectively store data at any scale

Kafka was designed to connect, store, and process data streams to power streaming data pipelines and applications across your business. As part of this, it's often desirable to store data in Kafka for event sourcing, caching, compliance, or other stream processing needs that require historical context. However, storing large amounts of data with open source Kafka can be tremendously expensive as it also requires growing compute, restricting how much these applications can leverage historical data.

With Confluent, we've decoupled storage from compute so you can store data at any scale as long as you need in a highly cost-effective manner:

- **Save on unnecessary compute** infrastructure as storage auto-scales 10x faster and easier
- **Never worry about capacity limits** with storage that grows automatically as you need it, without limits on retention time
- **Pay only for retained data**, not for any pre-provisioned or unused capacity

* Infinite Storage Enabled





02

Complete

How a complete and secure Kafka solution can improve time-to-value and TCO

RATHER THAN INVESTING significant engineering resources to build and maintain foundational tooling for Kafka, Confluent enables you to immediately stream data throughout your enterprise with a battle-tested, secure, and ready-to-use platform. With Confluent, you can deploy to production quickly and confidently while eliminating your team's day-to-day operational risks and burdens by leveraging a complete set of enterprise-grade features and services.

Confluent is so much more than Apache Kafka



Complete: Go beyond Kafka with all the essential tools for a complete data streaming platform



DEVELOPER

Unrestricted Developer Productivity



OPERATOR

Efficient Operations at Scale



ARCHITECT

Production-stage Prerequisites



EXECUTIVE

Partnership for Business Success

Multi-language Development
Stream Processing & Integration

Management & Monitoring
Flexible DevOps Automation

Enterprise-grade Security
Stream Governance

Complete Engagement Model
Business Case Justification

COMMITTER-DRIVEN EXPERTISE



Enterprise Support



Professional Services



Training



Partners



Unrestricted Developer Productivity

Confluent comes with a rich set of developer tools to boost engineering agility so that you can quickly build streaming applications & pipelines on top of Kafka to power rich, digital customer experiences and data-driven business operations:

- **Multi-language development:** Battle-tested clients for popular programming languages and REST proxy for any network-connected application, so that Kafka is broadly accessible to any developer
- **Stream processing and integration:** 120+ connectors for popular sources and sinks, SQL interface to enrich and process data streams for immediate action, and a visual builder to rapidly build, test, and deploy streaming data pipelines

Efficient Operations at Scale

Operating Kafka at scale can be immensely challenging, requiring manual, complex, and time-consuming management tasks to maintain the platform's performance, capacity, and availability. Confluent significantly reduces the operational burden of Kafka, whether you choose to offload management entirely to Confluent Cloud's fully managed service or self-manage with Confluent Platform:

- **Management and monitoring:** GUI tools to ensure cluster health, accelerate troubleshooting, and meet uptime SLAs
- **Flexible DevOps automation:** Infrastructure-as-code tooling to fully automate deployment and key lifecycle tasks

Production-Stage Prerequisites

Moving to production requires minimizing the risk of security breaches, data corruption, and downtime—concerns that only get harder and more complex as you scale. Confluent ensures you can meet stringent architectural requirements to avoid business disruption and any threats to your revenue-generating applications with a suite of enterprise-grade security and governance features:

- **Enterprise-grade security:** Role-based access control, audit logs for forensics and compliance, private networking, and encryption of sensitive data
- **Stream governance:** Discover, understand, and trust your data streams with Schema Registry, Stream Lineage, and Stream Catalog

Partnership for Business Success

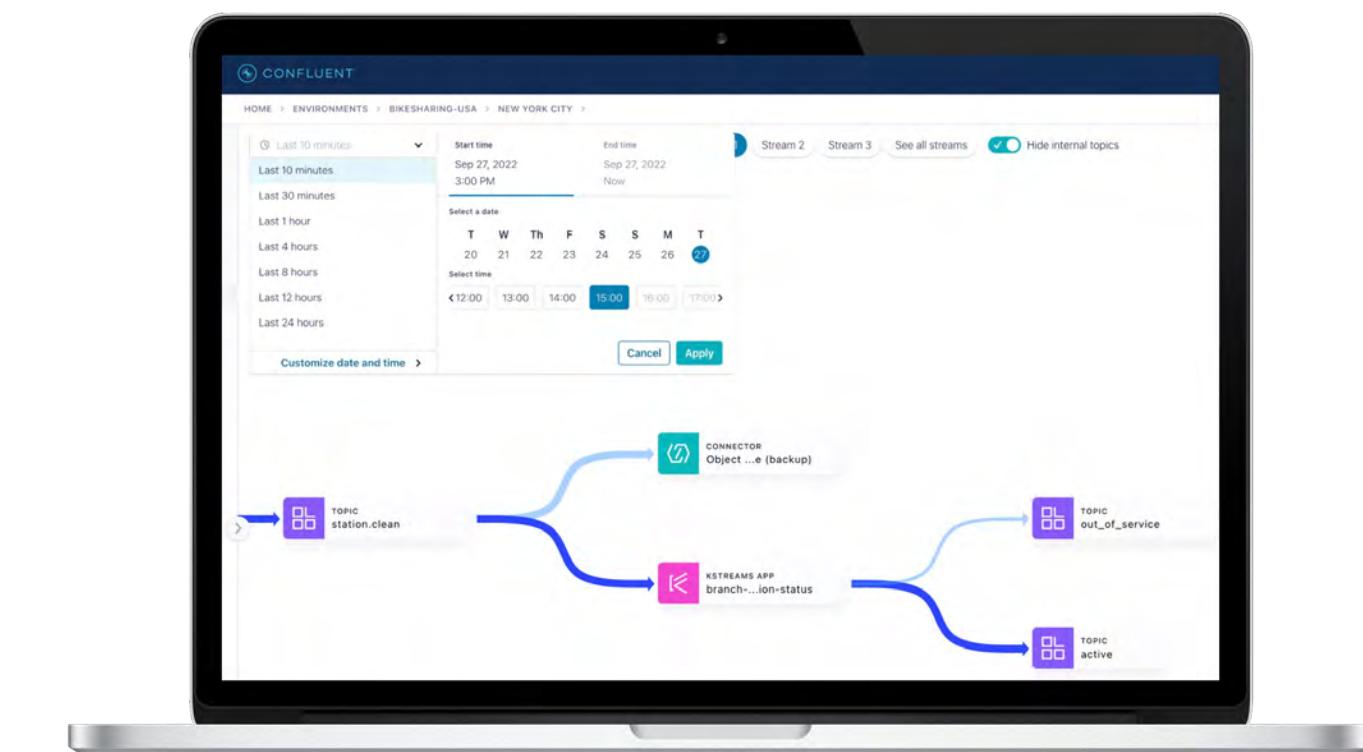
Confluent supports you throughout the software development lifecycle, ensuring your success for data streaming throughout your business:

- **Complete engagement model:** Partnership from your first data streaming project to building your central nervous system
- **Business case justification:** Reduce your Kafka TCO by up to 60% to improve the ROI and time-to-market of your data streaming projects

Committer-Driven Expertise

Confluent was founded by the original creators of Apache Kafka, providing unparalleled, industry-leading support and services from the team with 5M+ hours of experience. Ensuring our customers' success with data streaming is our singular focus:

- **Enterprise Support:** Access to expert support engineers that work directly with Apache Kafka committers
- **Professional Services:** Engagements to achieve your desired business outcomes and de-risk your investment in Kafka
- **Education:** Variety of training offerings to build broad Kafka expertise throughout your organization



03

Everywhere

Build hybrid and multi-cloud architectures with a platform that exists everywhere your data resides

KAFKA HELPS TO STREAMLINE digital architectures, avoid complex point-to-point integrations, and break data silos by syncing data systems and apps in real time. To truly meet that goal though, you need Kafka to span across all of your environments. If you operate across a multitude of on-premises and cloud environments, other data streaming solutions will only be able to support a subset of your overall architecture, limiting your ability to build a consistent data layer across your entire enterprise.





Deployment flexibility everywhere you need it

Confluent exists everywhere your applications and data reside, providing you the freedom to leverage Confluent Cloud—a fully managed service on the leading public clouds (AWS, Azure, Google Cloud) and Confluent Platform—self-managed software you can deploy for workloads that remain on-premises, whether on bare-metal, virtual machines, or Kubernetes. Confluent Cloud enables you to offload Kafka management entirely and instead focus on building applications that drive your business forward, helping to dramatically lower your Kafka TCO and accelerate time-to-value for new and existing use cases. Confluent Platform is designed to support workloads that need to remain on-prem, whether due to regulations or to cloud migration efforts that are more gradual.

Best of all, it's simple to connect and share data between independent Kafka clusters regardless of their deployment model, enabling you to bridge your entire data architecture across all of your different environments. When you're ready to move from on-prem to cloud and offload Kafka management entirely, Confluent provides Cluster Linking to easily migrate these self-managed clusters to our fully managed cloud service.

Confluent connects your data everywhere it resides



On-premises

Deploy on bare metal, VMs, or Kubernetes with **Confluent Platform**

Hybrid Cloud

Build a persistent bridge from datacenter to cloud with **Cluster Linking**

Public / Multi-Cloud

Leverage a fully managed service with **Confluent Cloud**



Customer Stories

How Instacart scales inventory management with Confluent's cloud-native platform

"We don't have the ability to go hire a bunch of engineers to babysit Kafka. Being able to offload those concerns to Confluent is just a major relief, allowing us to focus on delivering value to the organization and not worry about ops or other infrastructure overhead."

— Nate Kupp, Director of Engineering & Data, Instacart

Instacart experienced its highest customer demand in history at the onset of the COVID-19 pandemic. Its customer-facing applications, backend systems, and business processes were all put to the test as millions of people turned to Instacart as an essential service to safely get the groceries and goods they needed. Instacart required a platform that could support its end-to-end order fulfillment process from the time an item is added to a cart, to when it's delivered to the customer.

They turned to Confluent's cloud-native data streaming platform to connect the disparate systems and applications required to update inventory in real time. With the help of Confluent, Instacart was able to fully support and scale its underlying data infrastructure, enabling them to better reallocate valuable technical resources away from managing Kafka and instead toward innovating on core products for their customers.



How Michelin achieved 35% cost savings with real-time inventory management using Confluent Cloud

"We decided Confluent was necessary when we realized the limitations of self-managing open source Kafka on our own – the inability to scale cost-effectively, monitoring and security limitations, and no well-defined path to the cloud. Confluent offered a complete platform with all the enterprise-grade capabilities we needed to run mission-critical use cases, end to end. Over the last two years, Confluent has helped us to gain eight or nine months in terms of time-to-market to deploy the technology."

— Yves Caseau, Chief Digital & Information Officer, Michelin

Michelin relies heavily on data for all its manufacturing, services, and publishing lines of business. They had long dealt with unreliable and outdated reporting on inventory, which affected their global supply chain and logistics operations. Michelin began implementing Kafka in an on-prem model in its data centers, however soon realized they needed to migrate to the cloud in order to mitigate the operational challenges of managing Kafka.

To address these challenges and accelerate their journey to the cloud, Michelin chose Confluent Cloud to power all of their data streaming needs, significantly reducing operational costs by 35%. Confluent enabled Michelin to improve uptime with a 99.99% SLA, drive faster time-to-market with elastic scalability, and ultimately refocus its Kafka-dedicated teams to higher value projects.



Why you need Confluent

Every minute spent building out components needed to move Kafka into production or manually performing complex management tasks is time taken away from building rich customer experiences and improving your business efficiency with data-driven operations. Organizations that use Confluent have launched their streaming projects into production over six months faster than with Apache Kafka on their own, all while drastically reducing their day-to-day operational burdens, risks, and costs by up to 60%.

Confluent enables you to scale Kafka across your organization without the need to scale your teams.

Interested in learning more? [Contact us](#) today to see how Confluent can help ensure your success with Kafka and data streaming.

Is your team ready to try Confluent?
[Get started for free today.](#)

